

MuSHIN: A multi-way SMILES-based hypergraph inference network for metabolic model reconstruction

1

3 Yanlong Zhao^{1,†}, Yixiao Chen^{2,†}, Yi Yu³, Xiang Liu³, Jiawen Du⁴, Jun Wen⁵, Chen Liao⁶, Quan

4 Sun⁷, Ren Wang^{3,*}, Can Chen^{4,8,9,*}

5 ¹*Department of Electrical Computer Engineering, University of Rochester, Rochester, NY 14627, USA*

6 ²*Department of Computer Science, University of North Carolina at Chapel Hill, Chapel Hill, NC 27599, USA*

7 ³*Department of Electrical and Computer Engineering, Illinois Institute of Technology, Chicago, IL 60616, USA*

8 ⁴*Department of Biostatistics, University of North Carolina at Chapel Hill, Chapel Hill, NC 27599, USA*

9 ⁵*Harvard Medical School, Harvard University, Boston, MA 02115, USA*

10 ⁶*Department of Microbiology and Immunology, Dartmouth College, Hanover, NH 03755, USA*

11 ⁷*Center for Computational and Genomic Medicine, Children's Hospital of Philadelphia, Philadelphia, PA 19104, USA*

12 ⁸*School of Data Science and Society, University of North Carolina at Chapel Hill, Chapel Hill, NC 27599, USA*

13 ⁹*Department of Mathematics, University of North Carolina at Chapel Hill, Chapel Hill, NC 27599, USA*

14 [†] These authors contributed equally to this work.

15 * To whom correspondence should be addressed (rwang74@iit.edu; canc@unc.edu).

16 **Abstract**

17 Genome-scale metabolic models (GEMs) are indispensable tools for probing cellular metabolism,
18 enabling predictions of metabolic fluxes, guiding strain optimization, and advancing biomedical
19 research. However, their predictive capacity is often compromised by incomplete reaction
20 networks, stemming from gaps in biochemical knowledge, annotation inaccuracies, and insufficient
21 experimental validations. Here we present MuSHIN (**M**ulti-way **S**MILES-based **H**ypergraph
22 **I**nterface **Network), a novel deep hypergraph learning method that integrates network topology
23 with biochemical domain knowledge to predict missing reactions in GEMs. Evaluated on 926
24 high- and intermediate-quality GEMs with artificially removed reactions, MuSHIN significantly
25 outperforms state-of-the-art methods, achieving up to a 17% improvement across multiple metrics
26 and maintaining robust recovery even under severe network sparsity. Furthermore, MuSHIN
27 substantially enhances phenotypic predictions in 24 draft GEMs associated with fermentation by
28 resolving critical metabolic gaps, as validated against experimental measurements. Together, these
29 findings highlight MuSHIN's potential to advance GEM reconstruction and accelerate discoveries
30 in systems biology, metabolic engineering, and precision medicine.**

31 **Introduction**

32 Genome-scale metabolic models (GEMs) are foundational tools in systems biology, offering
33 a comprehensive representation of an organism's metabolic pathways at the genome scale^{1–5}.
34 By combining biochemical and physiological insights, GEMs allow researchers to simulate
35 cellular metabolism in silico and estimate metabolic fluxes under varying conditions. These
36 simulations inform strategies for genetic engineering and have driven advances in metabolic
37 engineering^{6,7}, drug discovery^{8,9}, and synthetic biology^{10,11}. Despite their broad utility, GEMs
38 are often incomplete due to gaps in biochemical knowledge, annotation errors, and limited
39 experimental validation^{12,13}. This incompleteness undermines the functional integrity of GEMs,
40 leading to inaccurate modeling of essential cellular phenotypes, such as biomass production,
41 metabolite secretion, and gene essentiality^{14–16}. The impact is particularly pronounced in non-
42 model organisms, whose metabolic pathways remain poorly characterized^{12,13}. Even in well-
43 studied species, certain reactions are only active under specific environmental or cellular
44 contexts^{17,18}. Addressing these deficiencies is thus critical for improving GEM fidelity and
45 expanding their applicability across organisms and experimental settings.

46 A variety of optimization-based strategies have been developed to address gaps in metabolic
47 models, many of which rely on incorporating phenotypic information, such as growth profiles
48 or flux distributions, to guide the insertion of missingg reactions^{14,19–23}. These methods seek to
49 reconcile model predictions with experimental observations by identifying inconsistencies and
50 supplementing pathways to restore metabolic function¹⁴. However, their reliance on condition-

51 specific experimental data restricts their use to well-characterized organisms, making them
52 unsuitable for many microbial species that are uncultivable or poorly annotated²⁴. In contrast,
53 topology-driven methods such as GapFind/GapFill²² offer a phenotype-independent alternative
54 by leveraging stoichiometric constraints to identify non-functional metabolites and proposing
55 reactions that enable feasible flux through the network (see Supplementary Note 1). Despite
56 circumventing the need for experimental measurements, these methods rely on linear optimization
57 and simplified network assumptions, overlooking the complex multi-metabolite interactions
58 fundamental to biochemical systems. As genome-scale models continue to expand in scope and
59 complexity, such simplifications increasingly constrain their ability to reflect biological reality.

60 Hyperedge prediction offers a promising solution to these challenges by leveraging
61 hypergraphs, a generalization of graphs in which hyperedges can simultaneously connect multiple
62 nodes^{25–30}. Metabolic networks can be naturally represented as hypergraphs, where each hyperedge
63 corresponds to a biochemical reaction involving multiple metabolites acting as substrates and
64 products^{26,30,31}. Unlike conventional graph-based methods that model pairwise interactions,
65 hypergraph-based representations preserve the higher-order connectivity inherent in metabolic
66 reactions, enabling a more accurate and comprehensive characterization of metabolic network
67 structures^{12,31–33}. However, existing hypergraph-based machine learning methods have notable
68 drawbacks (see Supplementary Note 1). For example, CMM³⁴ and C3MM³⁵ employ an integrated
69 training and prediction process that includes all candidate reactions from a predefined reaction
70 pool during training, limiting its scalability and requiring retraining whenever a new reaction pool
71 is introduced. The neural network-based method NHP³⁶ approximates hypergraphs using graphs

72 to generate node features, leading to the loss of higher-order interactions. Recent methods such
73 as CHESHIRE¹² and CLOSEgaps¹³ lack the expressive capacity of attention-based hypergraph
74 models, struggling to capture the complex multi-metabolite interactions inherent in metabolic
75 reactions. More importantly, these methods do not fully exploit accessible domain knowledge,
76 such as metabolite chemical structures, which are essential for evaluating reaction plausibility.
77 These shortcomings underscore the pressing need for biologically grounded, domain-informed
78 methods that preserve the hypergraph topology to more accurately recover missing reactions.

79 In this work, we introduce MuSHIN (**M**ulti-way **S**MILES-based **H**ypergraph **I**nterface
80 **N**etwork), a novel hypergraph-based deep learning method for predicting missing reactions
81 in GEMs by integrating both topological structure and biochemical semantics. MuSHIN
82 incorporates molecular-level information by embedding metabolite structures using transformer-
83 based models, enriching the hypergraph representation with chemical semantics. A dynamic
84 attention mechanism further refines metabolite and reaction features through iterative message
85 passing, allowing MuSHIN to effectively learn the complex higher-order interactions between
86 metabolites and reactions. Through comprehensive evaluations across multiple GEM databases,
87 MuSHIN significantly outperforms state-of-the-art methods in both internal validation based on the
88 recovery of synthetically removed reactions and external validation measured by the prediction of
89 experimentally observed phenotypes. By bridging biochemical knowledge with network topology,
90 MuSHIN offers a robust and scalable framework for automated GEM curation and advances
91 the predictive accuracy of metabolic modeling in systems biology, metabolic engineering, and
92 precision medicine.

93 **Results**

94 **A brief overview of MuSHIN.** MuSHIN is an advanced deep learning-based method designed to
95 predict missing reactions in GEMs by integrating both topological and biochemical semantics of
96 metabolic networks. It introduces two key innovations: (1) the use of transformer-based models,
97 including RXNFP³⁷ and ChemBERTa³⁸, to initialize chemically informed features of metabolites
98 and reactions; and (2) a dynamic attention mechanism³⁹ that iteratively refines metabolite and
99 reaction features through bidirectional node-hyperedge message passing. By jointly capturing
100 structural connectivity and chemical context, MuSHIN offers a powerful and generalizable
101 approach for reconstructing high-fidelity GEMs.

102 The metabolic network is represented as a hypergraph, where nodes denote metabolites
103 and hyperedges correspond to biochemical reactions (**Fig. 1a-b**; see “Hypergraphs” in Methods).
104 This formulation inherently captures the complex, higher-order interactions characteristic of
105 metabolic reactions, compared to conventional pairwise graphs. For model training, we curate
106 a balanced dataset comprising both positive and negative reactions, where positive samples
107 are drawn from well-established GEMs, while negative samples are synthetically generated by
108 perturbing known reactions to create chemically invalid alternatives, e.g., replacing one metabolite
109 with a non-reactive compound (**Fig. 1c**; see “Negative reaction generation” in Methods). This
110 deliberate generation strategy compels the model to discern biologically plausible reactions
111 from implausible ones by leveraging subtle structural and topological differences. All reactions,
112 including both positive and negative, are encoded into a hypergraph incidence matrix (**Fig. 1d**),

113 which captures metabolite-reaction connectivity as a binary relationship, forming the basis for
114 subsequent hypergraph learning. For simplicity, network directionality is not considered.

115 The architecture of MuSHIN consists of a feature initialization stage followed by iterative
116 feature refinement (**Fig. 1e**). During initialization, chemical information from metabolites and
117 reactions is encoded via transformer-based models applied to their SMILES representations, a
118 widely used textual notation that linearly encodes molecular structures. Specifically, metabolite
119 structures are embedded using ChemBERTa³⁸, while reaction SMILES are processed through
120 RXNFP³⁷, yielding rich high-dimensional feature vectors that capture detailed molecular semantics
121 (see “Metabolite and reaction embeddings” in Methods). These vectors are then projected
122 through multilayer perceptrons to a common dimensionality for subsequent integration. Following
123 initialization, MuSHIN refines these features using a hypergraph neural network enhanced
124 with a dual dynamic attention mechanism³⁹. This module iteratively updates metabolite (node)
125 and reaction (hyperedge) features via alternating node-to-hyperedge and hyperedge-to-node
126 attention-based message passing, allowing the model to integrate higher-order dependencies and
127 capture complex relationships within the metabolic network (see “Hypergraph neural network
128 architecture” in Methods). Unlike conventional methods that primarily depend on pooling
129 techniques, MuSHIN directly outputs refined context-aware reaction feature vectors, preserving
130 nuanced local and global structural information. In the final step, the refined reaction feature
131 vectors are fed into a single-layer neural network to predict the likelihood of reaction validity. The
132 model is trained end-to-end by minimizing the discrepancy between predicted scores and ground-
133 truth labels, enabling MuSHIN to faithfully capture the multifaceted structural and biochemical

¹³⁴ relationships inherent to GEMs (see Supplementary Note 2).

¹³⁵ **MuSHIN outperforms existing methods with artificially introduced gaps.** To assess the
¹³⁶ performance of MuSHIN in recovering missing reactions, we performed internal validation using
¹³⁷ GEMs with artificially introduced gaps (**Fig. 2a**, top testing scheme). Specifically, we evaluated
¹³⁸ MuSHIN against four representative baselines, including CLOSEgaps¹³, CHESHIRE¹², NHP³⁶,
¹³⁹ and HGNN²⁶ (see Supplementary Note 1), across 108 models from the BiGG database¹⁰ and
¹⁴⁰ 818 models from the AGORA database⁴⁰. Negative reactions were synthetically generated by
¹⁴¹ perturbing known reactions, ensuring chemical invalidity while preserving atom balance, and
¹⁴² paired with known valid reactions to create a balanced 1:1 positive-to-negative ratio. Each model
¹⁴³ was split into 60% training, 20% validation, and 20% testing sets. Details of the hyperparameter
¹⁴⁴ settings are provided in Supplementary Note 3. To ensure statistical robustness, we performed 10
¹⁴⁵ Monte Carlo runs per model and reported results as the average across runs.

¹⁴⁶ On 108 GEMs from the BiGG database, MuSHIN exhibits consistent superiority across
¹⁴⁷ multiple evaluation metrics when benchmarked against four baseline methods (**Fig. 2b-e**).
¹⁴⁸ Specifically, MuSHIN achieves a median F1 score of 93.69%, precision of 93.98%, recall of
¹⁴⁹ 93.49%, and an AUPRC of 95.20%. Relative to the strongest baseline, CLOSEgaps, MuSHIN
¹⁵⁰ improves the F1 score by 17.01% (P -value = 4.3E-25; paired t-test), precision by 16.65% (P -
¹⁵¹ value = 1.8E-18), recall by 12.86% (P -value = 1.6E-20), and AUPRC by 16.99% (P -value
¹⁵² = 2.1E-27). These improvements reflect significant enhancements in both sensitivity (recall)
¹⁵³ and specificity (precision), critical for accurate gap-filling in GEM reconstruction. Performance

154 margins over other baselines, such as CHESHIRE, NHP, and HGNN, are even more pronounced.
155 Beyond its strong average performance, MuSHIN exhibits lower performance variance across
156 GEMs of varying sizes, as indicated by tighter interquartile ranges, particularly in precision
157 and AUPRC (**Fig. 2c, e**). This consistency suggests that MuSHIN is robust across networks of
158 different complexity, in contrast to CLOSEgaps, which shows markedly higher variance, due to
159 its sensitivity to network size and training instability in smaller models. Furthermore, MuSHIN
160 maintains high performance under a range of evaluation settings, including different thresholding
161 criteria, negative sampling strategies, and various sampling ratios, consistently surpassing all
162 baselines (see Supplementary Note 3; Supplementary Figs. 1-5). Similar trends are observed on
163 the 818 GEMs from the AGORA database, where MuSHIN again outperforms all baselines across
164 all metrics (**Fig. 2f-i**). These results confirm that the advantages of MuSHIN are not confined to a
165 specific dataset but instead generalize robustly across diverse reconstruction platforms, organismal
166 systems, and metabolic network architectures, positioning MuSHIN as a state-of-the-art framework
167 for accurate, scalable, and reliable reaction recovery in genome-scale metabolic modeling.

168 **MuSHIN enables robust recovery in highly incomplete GEMs.** To assess the resilience of
169 MuSHIN in scenarios of severe network incompleteness, we conducted a second internal validation
170 experiment (**Fig. 2a**, bottom testing scheme). We selected nine representative GEMs from the
171 BiGG database, each comprising approximately 2,000 reactions, and systematically removed 20%,
172 40%, 60%, or 80% of their reactions to simulate missing data. The remaining reactions were
173 used for model training, supplemented with synthetically generated negative reactions at a 1:1
174 positive-to-negative ratio. For each level of degradation, we evaluated the model's ability to recover

175 missing reactions by measuring how many of the top N predicted reactions correctly matched to
176 the N removed ones. MuSHIN was benchmarked against the same four baselines, CLOSEgaps,
177 CHESHIRE, NHP, and HGNN, using identical model configurations and experimental protocols
178 as in the first type of internal validation. This setup provides a stringent test of the model's
179 recoverability under varying degrees of sparsity.

180 MuSHIN consistently achieves high recovery rates across the selected GEMs and varying
181 levels of missing rate, with an overall mean recovery approaching 90% (**Fig. 3**). When averaged
182 across all conditions, MuSHIN outperforms the strongest baseline, CLOSEgaps, by 8.74%,
183 demonstrating substantial and consistent gains over all competing methods. For example, in the
184 model STM_v1_0 (**Fig. 3h**), MuSHIN exceeds CLOSEgaps by roughly 10% across all tested
185 removal rates (P -value = 1.3E-4). At the particularly challenging 60% removal level, MuSHIN
186 maintains a recovery rate of 88.12%, whereas CLOSEgaps falls to 75.57%, underscoring the
187 superior resilience of MuSHIN under extreme data sparsity. Across most GEMs and removal
188 scenarios, the performance improvements of MuSHIN over baselines are statistically significant,
189 with paired P -values below 0.001 in the majority of cases, and below 0.01 in the remainder.
190 Additionally, the baseline methods exhibit notable variability in recovery performance as removal
191 rates increase. In particular, HGNN fails to generalize under high levels of network incompleteness,
192 often yielding recovery rates below 30% when 80% of reactions are withheld. In contrast,
193 MuSHIN remains remarkably stable across all levels of incompleteness and model sizes, reflecting
194 its capacity to generalize effectively even in data-scarce scenarios. These results highlight the
195 robustness of MuSHIN in recovering complex metabolic structure from limited information and

196 its suitability for reconstructing highly incomplete GEMs.

197 **MuSHIN advances accurate phenotypic prediction in fermentation.** To evaluate the biological
198 relevance and practical utility of MuSHIN, we performed external validation to test whether gap-
199 filling with MuSHIN enhances the prediction of actual metabolic phenotypes in draft GEMs
200 against experimental data¹². Our external validation workflow (**Fig. 4a**) begins by filtering
201 candidate reactions from a universal reaction pool (e.g., the universal BiGG reaction pool),
202 excluding those already present in the draft model or involving molecular oxygen. MuSHIN
203 then scores the remaining candidates using a composite metric that integrates multiple evidence
204 sources, including model-based scores and similarity coefficients, to prioritize reactions for gap-
205 filling. We subsequently applied flux balance analysis to both the original draft and MuSHIN-
206 augmented models to predict metabolic phenotypes. These predictions were rigorously compared
207 to experimental phenotypic data using established classification metrics, enabling quantitative
208 evaluation of MuSHIN’s contribution to improving metabolic phenotype prediction. Detailed steps,
209 including similarity coefficients and flux balance analysis, can be found in Supplementary Note 5.

210 We applied MuSHIN to a set of 24 draft GEMs (Supplementary Table 1) derived from
211 anaerobic bacterial species⁴¹ and evaluated its ability to improve phenotypic predictions of
212 fermentation (see Supplementary Note 5). Specifically, we incorporated top 100 reactions
213 prioritized by MuSHIN (MuSHIN-100) into each draft model and benchmarked the resulting
214 performance against three comparators: the unmodified draft models (BEFORE), CHESHIRE-
215 augmented models (CHESHIRE-100)¹², and CLOSEgaps-augmented models (CLOSEgaps-

216 100)¹³. As shown in **Fig. 4b-e**, MuSHIN-augmented models exhibit substantial improvements
217 across multiple evaluation metrics, including F1 score, precision, recall, and AUPRC, compared to
218 all the other models. MuSHIN achieves median values exceeding 40% for F1 score, precision, and
219 recall, whereas the corresponding metrics for all other models remain zero (P -values = 7.3E-3,
220 2.1E-2, and 7.0E-3, respectively; Wilcoxon signed-rank test), underscoring the unique advantages
221 conferred by MuSHIN in recovering biologically meaningful phenotypic predictions. Notably,
222 previous studies have shown that both CHESHIRE and CLOSEgaps require the addition of at least
223 200 reactions to observe any significant improvement in predictive performance^{12,13}. A detailed
224 comparison of fermentation product predictions between draft and MuSHIN-augmented models is
225 shown in **Fig. 4f**, where MuSHIN successfully corrects several false-negative predictions from
226 the draft GEMs, enhancing agreement with experimentally observed fermentation phenotypes.
227 Additionally, MuSHIN consistently outperforms CHESHIRE, CLOSEgaps, and the non-specific
228 addition of 100 universal reactions (Supplementary Fig. 7). These findings highlight MuSHIN's
229 predictive strength and practical utility in improving the functional fidelity of GEMs under
230 experimentally grounded and biologically realistic conditions, making it a potential tool for
231 advancing biotechnology, microbial ecology, and systems biology research.

232 **MuSHIN resolves critical gaps in GEMs.** To further understand the underlying biological
233 advances provided by MuSHIN, we performed enrichment analyses for reactions prioritized
234 by MuSHIN in the previous external experiment. Our results reveal a striking enrichment of
235 transporter and exchange reactions. Among unique reactions proposed by MuSHIN (e.g., in the
236 MuSHIN-100 set), 49.92% are annotated as transport or exchange, representing a significant

enrichment compared to their prevalence in universal reaction databases (28.1% in the BiGG universal reaction pool; P -value < 1.0E-50; Fisher's exact test). This selective enrichment aligns with the "membrane bottleneck hypothesis" in metabolic engineering, where transport limitations often constrain metabolic flux more than enzymatic capacity, indicating that MuSHIN selectively identifies boundary reactions in a non-random, biologically grounded manner.

To elucidate the mechanistic basis for MuSHIN's improved phenotypic predictions, we compared its performance against the CHESHIRE-100 baseline across 24 selected genomes (Supplementary Table 1). MuSHIN improves the phenotypic prediction F1 score in 11 of these genomes, yielding substantial gains in forecasting specific metabolic products. Notably, MuSHIN enables the correct prediction of formic acid in five genomes and acetic acid in five other genomes where CHESHIRE-100 fails. Additional significant improvements are observed for ethanol, which is uniquely predicted by MuSHIN in three distinct genomes missed by CHESHIRE. Furthermore, DL-lactic acid, succinic acid, and propionic acid are each correctly identified in one additional genome. Collectively, MuSHIN introduces 16 new correct product predictions absent in the CHESHIRE-100 results, highlighting its superior capacity to recover missing metabolic functions. Further analysis (Supplementary Table 2) reveals that these improvements are likely explained by MuSHIN's addition of specific transporter families, which enhance substrate secretion or cofactor availability. Key examples include transporters for formate (FORt2*, FORt3*), CO₂/O₂ exchange (CO2tex, O2tex), Zn²⁺ (ZN2t*, Zn2tex), ethanol (ETOHt*), and fumarate/malate (FUMt*, MALt*). For instance, improved formic acid secretion is largely due to the addition of specific formate transporters (FORt2*, FORt3*) in six genomes⁴², often coupled with CO₂

258 exchange transporters (CO2tex) that facilitate carbon metabolism⁴³. Similarly, enhanced ethanol
259 production is led by the addition of ethanol-specific transporters (ETOHt2r*, ETOHt3*) in five
260 genomes, enabling efficient product efflux⁴⁴.

261 Beyond single-transporter corrections, MuSHIN uniquely resolves complex metabolic
262 gaps. A notable instance in *GCF_000144405.1* involves the restoration of succinate production
263 through coordinated addition of fumarate transporters (FUMt2_3pp, FUMt2_2pp, FUMtpp),
264 facilitating the reductive branch of the tricarboxylic acid cycle (TCA) under anaerobic conditions⁴⁵.
265 Similarly, in *GCF_000013285.1*, improvements in both ethanol and formic acid production are
266 achieved through the addition of ethanol-specific transporters (ETOHt2rpp, ETOHt3, ETOHt2r)
267 and formate transporters (FORt2pp, FORt3, FORt2)⁴⁶, demonstrating MuSHIN's ability to
268 enhance multiple fermentation pathways simultaneously. In genomes like *GCF_000392875.1* and
269 *GCF_000469345.1*, MuSHIN enhances ethanol production even without the addition of dedicated
270 ethanol transporters, suggesting the involvement of alternative transport mechanisms or metabolic
271 rerouting. This synergy highlights MuSHIN's strength in uncovering both direct transporter-
272 mediated solutions and indirect metabolic enhancements—complex scenarios that simpler gap-
273 filling methods often miss. Notably, zinc transporters (ZN2t*, Zn2tex) are added across all eleven
274 genomes, indicating their fundamental role as cofactor suppliers for enhanced enzymatic activity in
275 metabolic pathway restoration. The universal requirement for zinc transporters reflects the metal's
276 role as a cofactor in over 300 enzymes, particularly those involved in central carbon metabolism
277 and redox reactions⁴⁵.

278 In summary, these findings demonstrate that MuSHIN's superior performance arises
279 from addressing a fundamental limitation in current gap-filling approaches: the systematic
280 underestimation of membrane transport constraints. By prioritizing functionally critical transport
281 and exchange reactions that alleviate these bottlenecks, and resolving complex multi-reaction gaps
282 through coordinated transporter addition, MuSHIN reinstates complete metabolic branches while
283 maintaining the stoichiometric and thermodynamic constraints essential for biological fidelity.

284 **Discussion**

285 In this work, we introduced MuSHIN (**M**ulti-way **S**MILES-based **H**ypergraph **I**nterface **Network),
286 a deep hypergraph learning framework that integrates metabolic network topology with
287 biochemical domain knowledge to address the persistent challenge of reaction incompleteness
288 in GEMs. By representing metabolic networks as hypergraphs, MuSHIN captures the intrinsic
289 high-order relationships among metabolites and reactions, offering a more faithful abstraction
290 of biochemical systems than traditional pairwise graph representations. In addition, MuSHIN
291 incorporates chemically informed molecular embeddings derived from SMILES strings, enabling
292 the model to reason about molecular structure and reactivity. A dynamic attention mechanism
293 further enhances this process by selectively emphasizing the most informative topological and
294 biochemical features, allowing MuSHIN to adaptively prioritize relevant signals during training
295 and prediction. The contributions of these components are further supported by ablation results
296 (Supplementary Fig. 6). Compared to recent state-of-the-art methods such as CHESHIRE
297 and CLOSEgaps, which primarily rely on reaction co-occurrence patterns or curated reaction**

298 similarity metrics, MuSHIN offers a fundamentally different learning paradigm, which allows
299 MuSHIN to generalize more effectively, particularly in sparse or incompletely annotated networks,
300 and establishes a robust foundation for scalable and biologically grounded metabolic network
301 reconstruction.

302 The effectiveness of MuSHIN was rigorously evaluated through both internal and external
303 validation. Internally, we conducted extensive benchmarking on a large collection of 926
304 high- and intermediate-quality GEMs, where reactions were systematically and artificially
305 removed to simulate real-world incompleteness. Across multiple quantitative metrics, MuSHIN
306 consistently outperforms existing state-of-the-art hypergraph-based gap-filling methods such
307 as CHESHIRE and CLOSEgaps. Notably, MuSHIN maintains strong predictive performance
308 even under conditions of extreme network sparsity, demonstrating its ability to infer missing
309 reactions in challenging scenarios where conventional approaches falter. Externally, we applied
310 MuSHIN to 24 draft GEMs focused on microbial fermentation pathways and validated the
311 predicted reaction additions by comparing model-derived phenotypic outputs with independent
312 experimental measurements. Incorporation of 100 reactions predicted by MuSHIN leads to
313 significant improvements in fermentation phenotype predictions, whereas other methods such
314 as CHESHIRE and CLOSEgaps fail to produce noticeable enhancements. This demonstrates
315 MuSHIN's superior ability to identify functionally relevant reactions that directly translate to more
316 accurate and biologically meaningful phenotypic outcomes.

317 Enhancing the completeness and fidelity of GEMs holds transformative potential across

318 diverse domains of biology, biotechnology, and medicine. In biotechnology, high-quality GEMs
319 enable rational strain engineering by revealing metabolic bottlenecks and guiding the optimization
320 of biosynthetic pathways for the production of pharmaceuticals, biofuels, and other high-value
321 compounds. In the context of human health, GEMs serve as powerful tools for precision medicine,
322 allowing the modeling of tissue- or patient-specific metabolism to uncover disease mechanisms,
323 identify biomarkers, and inform therapeutic interventions. In microbial ecology and synthetic
324 biology, more comprehensive metabolic reconstructions facilitate the design of synthetic consortia
325 and enable the strategic allocation of metabolic functions across community members. By
326 advancing automated, data-driven approaches for metabolic gap-filling such as MuSHIN, we not
327 only improve the structural and functional accuracy of GEMs but also unlock their full potential as
328 predictive frameworks for understanding and engineering complex biological systems.

329 Building on MuSHIN's strong performance, several avenues offer promising opportunities
330 for further development. First, while the framework effectively integrates molecular structure and
331 metabolic network topology, it currently lacks integration of additional biological layers, such
332 as gene-protein-reaction associations and regulatory interactions, which could be harnessed to
333 construct more comprehensive, multi-modal knowledge hypergraphs. Second, although SMILES-
334 based embeddings capture chemical properties, they may not fully reflect enzyme-specific context
335 or regulatory constraints, which are critical for accurate functional interpretation. Third, MuSHIN
336 currently treats gap-filling as a static prediction task and does not incorporate dynamic cellular
337 states or context-dependent metabolic shifts, which are increasingly important for modeling
338 complex biological systems. Pursuing these directions will support broader applicability of

339 MuSHIN and further advance the automation and accuracy of GEM reconstruction.

340 **Methods**

341 **Hypergraphs.** Hypergraphs extend traditional graphs by incorporating hyperedges (also known as
342 hyperlinks) that can connect any number of nodes³⁰. This flexibility allows hypergraphs to model
343 complex correlations in real-world data that go beyond simple pairwise interactions³⁷. Formally,
344 an unweighted hypergraph $\mathcal{H} = \{\mathcal{V}, \mathcal{E}\}$ comprises a set of nodes $\mathcal{V} = \{v_1, v_2, \dots, v_n\}$ and a set
345 of hyperedges $\mathcal{E} = \{e_1, e_2, \dots, e_m\}$, where each hyperedge $e_p \subseteq \mathcal{V}$ for $p = 1, 2, \dots, m$. Two
346 nodes are considered adjacent if they share a common hyperedge, and a hypergraph is connected if
347 there exists a path between any two nodes via a sequence of hyperedges. The incidence matrix of a
348 hypergraph, denoted by $\mathbf{H} \in \{0, 1\}^{n \times m}$, captures the relationships between nodes and hyperedges.
349 Specifically, the entry \mathbf{H}_{ip} is set to one if node v_i is part of hyperedge e_p , and zero otherwise.

350 **Negative reaction generation.** Predicting missing reactions in metabolic networks necessitates a
351 dataset containing both positive (existing) and negative (non-existing) reactions. We developed a
352 negative sampling strategy that generates invalid reactions by perturbing known reactions. Given a
353 set of known reactions $\mathcal{R}^+ = \{r_1, r_2, \dots, r_N\}$, where each reaction r_i comprises specific reactants
354 and products, we generate negative reactions through the following steps: (i) aggregate a metabolite
355 pool \mathcal{M} by combining metabolites from \mathcal{R}^+ with those from external databases; (ii) for each
356 reaction $r_i \in \mathcal{R}^+$, randomly select a fraction α of its metabolites for replacement; (iii) substitute
357 the selected metabolites with randomly chosen ones from $\mathcal{M} \setminus r_i$ (i.e., metabolites not present in

358 $r_i)$; (iv) ensure that the newly generated reaction r_i^- does not exist in \mathcal{R}^+ .

359 This procedure yields a set of negative reactions $\mathcal{R}^- = \{r_1^-, r_2^-, \dots, r_N^-\}$ that are
360 structurally similar to positive reactions but are chemically invalid due to improbable metabolite
361 combinations^{47,48}. Each reaction is represented by a stoichiometric vector $s_i \in \mathbb{R}^M$, where M
362 is the total number of unique metabolites, and each component $s_{i,j}$ represents the stoichiometric
363 coefficient n_j of metabolite j :

$$s_{i,j} = \begin{cases} -n_j, & \text{if metabolite } j \text{ is a reactant in } r_i, \\ n_j, & \text{if metabolite } j \text{ is a product in } r_i, \\ 0, & \text{otherwise.} \end{cases}$$

364 The complete dataset is $\mathcal{R} = \mathcal{R}^+ \cup \mathcal{R}^-$, with each reaction labeled as follows:

$$y_i = \begin{cases} +1, & \text{if } r_i \in \mathcal{R}^+, \\ -1, & \text{if } r_i \in \mathcal{R}^-. \end{cases}$$

365 **Reaction embedding using RXNFP.** RXNFP³⁷ is a transformer-based model pre-trained on
366 chemical reaction data, designed to encode reaction SMILES strings into numerical representations
367 that capture underlying chemical transformations and mechanistic patterns. Each reaction SMILES
368 string is tokenized and processed through multiple transformer layers, producing a fixed-length
369 reaction fingerprint (typically 256-dimensional) suitable for downstream tasks such as reaction
370 classification and pathway prediction. Mathematically, the RXNFP model processes a tokenized

371 reaction SMILES sequence $r = [r_1, r_2, \dots, r_N]$ through transformer layers as follows:

$$\mathbf{H}^{(l)} = \text{TransformerLayer}^{(l)}(\mathbf{H}^{(l-1)}),$$

372 where $\mathbf{H}^{(0)}$ represents the initial token embeddings, and $\mathbf{X}^{(l)} \in \mathbb{R}^{N \times d}$ denotes the hidden states

373 at layer l , with d being the hidden size. The final reaction fingerprint $\mathbf{f} \in \mathbb{R}^d$ is obtained by

374 applying a pooling operation over the last layer's hidden states $\mathbf{f} = \text{Pooling}(\mathbf{H}^{(L)})$. By integrating

375 RXNFP-based embeddings, MuSHIN effectively incorporates chemical reaction patterns into the

376 hypergraph model, improving its ability to predict missing reactions.

377 **Metabolite embedding using ChemBERTa.** Similarly, ChemBERTa³⁸ is a transformer-based

378 language model pre-trained on molecular data, which encodes metabolite SMILES strings

379 into meaningful vector representations. ChemBERTa captures complex structural and chemical

380 relationships, providing robust numerical features for metabolic network analysis. Metabolite data

381 is sourced from the BiGG database¹⁰, and corresponding SMILES representations are retrieved

382 via ChEBI IDs⁴⁷. If unavailable, external databases such as KEGG are consulted. Each SMILES

383 string is tokenized and processed through multiple transformer layers, producing a fixed-length

384 (768-dimensional) vector representation. Mathematically, for a tokenized SMILES sequence $s =$

385 $[s_1, s_2, \dots, s_N]$, the ChemBERTa model generates hidden states at each transformer layer. The

386 hidden states from the final layer L are denoted as

$$\mathbf{H}^{(L)} = \begin{bmatrix} \mathbf{h}_1^{(L)} \\ \mathbf{h}_2^{(L)} \\ \vdots \\ \mathbf{h}_N^{(L)} \end{bmatrix} \in \mathbb{R}^{N \times d},$$

387 where $\mathbf{h}_i^{(L)} \in \mathbb{R}^d$ represents the hidden state of the i th token, and $d = 768$. The metabolite's
388 vector representation $\mathbf{v} \in \mathbb{R}^d$ is computed by averaging these hidden states $\mathbf{v} = \frac{1}{N} \sum_{i=1}^N \mathbf{h}_i^{(L)}$.

389 By leveraging ChemBERTa-based embeddings, MuSHIN incorporates rich molecular properties
390 into its hypergraph representation, allowing for a more accurate characterization of metabolite
391 relationships within metabolic networks. The combination of RXNFP and ChemBERTa ensures
392 that MuSHIN effectively captures both topological and chemical features, enabling superior
393 performance in predicting missing reactions.

394 **Hypergraph neural network architecture.** Various hypergraph neural network models have been
395 proposed^{27,33,49–52}. However, most of these approaches simplify the problem by decomposing
396 hyperedges into pairwise interactions and then applying standard graph neural networks, which can
397 lead to a loss of high-order structural information. To effectively model the complex interactions
398 within the metabolic network, we propose a novel hypergraph neural network with a dual attention
399 mechanism that leverages both node and hyperedge features. Our architecture integrates advanced
400 attention mechanisms to capture higher-order dependencies and interactions within the hypergraph.
401 Specifically, the model iteratively refines feature representations through a combination of Node-
402 to-Edge and Edge-to-Node attention mechanisms, allowing MuSHIN to dynamically learn the

403 relationships between metabolites and reactions.

404 The Node-to-Edge attention mechanism aggregates information from nodes to hyperedges,
405 updating hyperedge features based on the features of their incident nodes. This enables hyperedges
406 to focus on the most relevant information from connected nodes, weighted by learned attention
407 coefficients. Given node features $\mathbf{X} \in \mathbb{R}^{N \times F_n}$ and the incidence matrix $\mathbf{H} \in \{0, 1\}^{N \times M}$, where
408 N is the number of nodes, M is the number of hyperedges, and F_n is the dimensionality of node
409 features, we first transform the node features into a shared latent space

$$\mathbf{Z}_n = \sigma(\mathbf{X}\mathbf{W}_n),$$

410 where $\mathbf{W}_n \in \mathbb{R}^{F_n \times F_h}$ is a learnable weight matrix, F_h is the hidden dimension, and σ is an
411 activation function (e.g., ReLU). For each hyperedge e , we compute attention coefficients for its
412 incident nodes

$$\alpha_{ie} = \frac{\exp(\text{LeakyReLU}(\mathbf{a}^\top [\mathbf{z}_i \| \mathbf{c}_e]))}{\sum_{v_j \in \mathcal{V}_e} \exp(\text{LeakyReLU}(\mathbf{a}^\top [\mathbf{z}_j \| \mathbf{c}_e])),$$

413 where \mathbf{z}_i is the transformed feature of node v_i , $\mathbf{a} \in \mathbb{R}^{2F_h}$ is a learnable attention vector, $\|$ denotes
414 vector concatenation, $\mathbf{c}_e \in \mathbb{R}^{F_h}$ is a learnable hyperedge representation (initialized randomly or as
415 zeros), and \mathcal{V}_e denotes the set of nodes connected to hyperedge e . The hyperedge feature \mathbf{h}_e is then
416 updated by aggregating the node features weighted by the attention coefficients $\mathbf{h}_e = \sum_{v_i \in \mathcal{V}_e} \alpha_{ie} \mathbf{z}_i$.
417 This mechanism allows hyperedges to emphasize the most informative nodes, enhancing their
418 ability to represent complex relationships within the hypergraph.

419 The Edge-to-Node attention mechanism updates node features by aggregating information
420 from the hyperedges they are connected to. By attending to the most relevant hyperedges, nodes

421 can incorporate higher-order relational information into their representations. First, we transform
 422 both node and hyperedge features into a shared latent space

$$\mathbf{Z}_n = \sigma(\mathbf{X}\mathbf{W}_n), \quad \mathbf{Z}_e = \sigma(\mathbf{H}^\top \mathbf{h}_e \mathbf{W}_e),$$

423 where $\mathbf{W}_e \in \mathbb{R}^{F_h \times F_h}$ is a learnable weight matrix for hyperedges. For each node i , the attention
 424 coefficients for connected hyperedges are computed as

$$\beta_{ie} = \frac{\exp(\text{LeakyReLU}(\mathbf{b}^\top [\mathbf{z}_i \| \mathbf{z}_e]))}{\sum_{k \in \mathcal{E}_i} \exp(\text{LeakyReLU}(\mathbf{b}^\top [\mathbf{z}_i \| \mathbf{z}_k])),$$

425 where \mathbf{z}_e is the transformed feature of hyperedge e , $\mathbf{b} \in \mathbb{R}^{2F_h}$ is a learnable attention vector, and
 426 \mathcal{E}_i denotes the set of hyperedges connected to node i . The updated node feature \mathbf{x}'_i is obtained by
 427 aggregating the hyperedge features weighted by the attention coefficients and adding a residual
 428 connection $\mathbf{x}'_i = \sigma(\sum_{e \in \mathcal{E}_i} \beta_{ie} \mathbf{z}_e) + \mathbf{z}_i$. This approach allows nodes to integrate contextual
 429 information from relevant hyperedges, resulting in richer feature representations.

430 We iteratively apply the Node-to-Edge and Edge-to-Node attention mechanisms to
 431 progressively refine both node and hyperedge features. At each iteration, updated node features
 432 inform the hyperedge features, and updated hyperedge features subsequently enhance the node
 433 features. This iterative process facilitates deep integration of structural information across the
 434 hypergraph, capturing complex dependencies and interactions. The iterative refinement process
 435 is formalized as follows:

$$\mathbf{H}^{(l)} = \text{NodeToEdgeAttention}(\mathbf{Z}_n^{(l-1)}, \mathbf{C}_e^{(l-1)}),$$

$$\mathbf{Z}_n^{(l)} = \text{EdgeToNodeAttention}(\mathbf{Z}_n^{(l-1)}, \mathbf{H}^{(l)}, \mathbf{C}_e^{(l-1)}),$$

436 for $l = 1, 2, \dots, L$, where $\mathbf{Z}_n^{(0)} = \sigma(\mathbf{X}\mathbf{W}_n)$ are the initial node features, $\mathbf{H}^{(l)}$ represents the
437 hyperedge features at layer l , and $\mathbf{C}_e^{(l-1)}$ are the hyperedge representations from the previous layer.
438 This iterative refinement enhances the model's ability to capture higher-order dependencies and
439 interactions within the hypergraph, leading to improved performance on downstream tasks such
440 as reaction prediction and network analysis (see Supplementary Note 3 for a detailed prediction
441 workflow).

442 Data availability

443 The datasets used and analyzed during the current study are included within this article
444 and its supplementary information files. The raw data were collected from publicly available
445 databases: ChEBI (<https://www.ebi.ac.uk/chebi/>), BiGG Models ([http://bigg.
446 ucsd.edu/](http://bigg.ucsd.edu/)), AGORA Models (<https://www.vmh.life>). More details can be found in
447 Supplementary Note 5.

448 Code availability

449 The source code for our framework is available at Github [<https://github.com/cyixiao/MuSHIN>].

450 References

- 452 1. Ines Thiele, Nathan D Price, Thuy D Vo, and Bernhard Ø Palsson. Candidate metabolic
453 network states in human mitochondria. *FEBS Letters*, 579(30):6435–6441, 2005.

- 454 2. Ines Thiele, Neema Jamshidi, Ronan MT Fleming, and Bernhard Ø Palsson. Genome-scale
455 reconstruction of escherichia coli's transcriptional and translational machinery: A knowledge
456 base, its mathematical formulation, and its functional characterization. *PLoS Computational
457 Biology*, 5(3):e1000312, 2009.
- 458 3. Sang Yup Lee and Hyun Uk Kim. Systems strategies for developing industrial microbial
459 strains. *Nature Biotechnology*, 33(10):1061–1072, Oct 2015.
- 460 4. Changdai Gu, Gi Bae Kim, Won Jun Kim, Hyun Uk Kim, and Sang Yup Lee. Current status
461 and applications of genome-scale metabolic models. *Genome biology*, 20:1–18, 2019.
- 462 5. Christian Lieven, Moritz E Beber, Brett G Olivier, Frank T Bergmann, Meric Ataman, Parizad
463 Babaei, Jennifer A Bartell, Lars M Blank, Siddharth Chauhan, Kevin Correia, et al. Memote
464 for standardized genome-scale metabolic model testing. *Nature biotechnology*, 38(3):272–
465 276, 2020.
- 466 6. Evangelos Simeonidis and Nathan D Price. Genome-scale modeling for metabolic
467 engineering. *Journal of Industrial Microbiology and Biotechnology*, 42(3):327–338, 2015.
- 468 7. Byoungjin Kim, Won Jun Kim, Dong In Kim, and Sang Yup Lee. Applications of genome-
469 scale metabolic network model in metabolic engineering. *Journal of industrial microbiology
470 and biotechnology*, 42(3):339–348, 2015.
- 471 8. Vytautas Raškevičius, Valeryia Mikalayeva, Ieva Antanavičiūtė, Ieva Ceslevičienė,
472 Vytenis Arvydas Skeberdis, Visvaldas Kairys, and Sergio Bordel. Genome scale metabolic
473 models as tools for drug design and personalized medicine. *PloS one*, 13(1):e0190636, 2018.

- 474 9. Jonathan L Robinson and Jens Nielsen. Anticancer drug discovery through genome-scale
475 metabolic modeling. *Current Opinion in Systems Biology*, 4:1–8, 2017.
- 476 10. Zachary A King, Justin Lu, Andreas Dräger, Philip Miller, Stephen Federowicz, Joshua A
477 Lerman, Ali Ebrahim, Bernhard O Palsson, and Nathan E Lewis. Bigg models: A platform
478 for integrating, standardizing and sharing genome-scale models. *Nucleic acids research*,
479 44(D1):D515–D522, 2016.
- 480 11. Jens Nielsen and Jay D. Keasling. Engineering cellular metabolism. *Cell*, 164(6):1185–1197,
481 2016.
- 482 12. Can Chen, Chen Liao, and Yang-Yu Liu. Teasing out missing reactions in genome-scale
483 metabolic networks through hypergraph learning. *Nature Communications*, 14(1):2375, 2023.
- 484 13. Xiaoyi Liu, Hongpeng Yang, Chengwei Ai, Ruihan Dong, Yijie Ding, Qianqian Yuan, Jijun
485 Tang, and Fei Guo. A generalizable framework for unlocking missing reactions in genome-
486 scale metabolic networks using deep learning. *arXiv preprint arXiv:2409.13259*, 2024.
- 487 14. Shu Pan and Jennifer L Reed. Advances in gap-filling genome-scale metabolic models
488 and model-driven experiments lead to novel metabolic discoveries. *Current opinion in
489 biotechnology*, 51:103–108, 2018.
- 490 15. Matthew N Benedict, Michael B Mundy, Christopher S Henry, Nicholas Chia, and Nathan D
491 Price. Likelihood-based gene annotations for gap filling and quality assessment in genome-
492 scale metabolic models. *PLoS computational biology*, 10(10):e1003882, 2014.

- 493 16. Peter D Karp, Daniel Weaver, and Mario Latendresse. How accurate is automated gap filling
494 of metabolic models? *BMC systems biology*, 12:1–11, 2018.
- 495 17. Daniel Machado, Sergej Andrejev, Michele Tramontano, and Kiran R Patil. Fast automated
496 reconstruction of genome-scale metabolic models for microbial species and communities.
497 *Nucleic acids research*, 46(15):7542–7553, 2018.
- 498 18. Changdai Gu, Gi Bae Kim, Won Jun Kim, Hyun Uk Kim, and Sang Yup Lee. Current status
499 and applications of genome-scale metabolic models. *Genome Biology*, 20(1):121, June 2019.
- 500 19. Jeffrey D Orth and Bernhard Ø Palsson. Systematizing the generation of missing metabolic
501 knowledge. *Biotechnology and bioengineering*, 107(3):403–412, 2010.
- 502 20. Wheaton L Schroeder and Rajib Saha. Optfill: a tool for infeasible cycle-free gapfilling of
503 stoichiometric metabolic models. *IScience*, 23(1), 2020.
- 504 21. Sylvain Prigent, Clémence Frioux, Simon M Dittami, Sven Thiele, Abdelhalim Larhlimi,
505 Guillaume Collet, Fabien Gutknecht, Jeanne Got, Damien Eveillard, Jérémie Bourdon, et al.
506 Meneco, a topology-based gap-filling tool applicable to degraded genome-wide metabolic
507 networks. *PLoS computational biology*, 13(1):e1005276, 2017.
- 508 22. Vinay Satish Kumar, Madhukar S Dasika, and Costas D Maranas. Optimization based
509 automated curation of metabolic reconstructions. *BMC bioinformatics*, 8:1–16, 2007.
- 510 23. Christopher S Henry, Matthew DeJongh, Aaron A Best, Paul M Frybarger, Ben Lindsay, and
511 Rick L Stevens. High-throughput generation, optimization and analysis of genome-scale
512 metabolic models. *Nature biotechnology*, 28(9):977–982, 2010.

- 513 24. Alexandre Almeida, Alex L Mitchell, Miguel Boland, Samuel C Forster, Gregory B Gloor,
514 Aleksandra Tarkowska, Trevor D Lawley, and Robert D Finn. A new genomic blueprint of the
515 human gut microbiota. *Nature*, 568(7753):499–504, 2019.
- 516 25. Can Chen and Yang-Yu Liu. A survey on hyperlink prediction. *IEEE Transactions on Neural*
517 *Networks and Learning Systems*, 35(11):15034–15050, 2024.
- 518 26. Yifan Feng, Haoxiang You, Zhirong Zhang, Rongrong Ji, and Yue Gao. Hypergraph neural
519 networks. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 33, pages
520 3558–3565, 2019.
- 521 27. Yiwei Bai, Hao Ding, Yifan Sun, Wenyu Wang, and Yihong Gong. Hypergraph convolution
522 and hypergraph attention. *Pattern Recognition*, 110:107637, 2021.
- 523 28. Can Chen, Amit Surana, Anthony M Bloch, and Indika Rajapakse. Controllability of
524 hypergraphs. *IEEE Transactions on Network Science and Engineering*, 8(2):1646–1657, 2021.
- 525 29. Can Chen and Indika Rajapakse. Tensor entropy for uniform hypergraphs. *IEEE Transactions*
526 *on Network Science and Engineering*, 7(4):2889–2900, 2020.
- 527 30. Claude Berge. *Hypergraphs: combinatorics of finite sets*, volume 45. Elsevier, 1984.
- 528 31. Dengyong Zhou, Jiayuan Huang, and Bernhard Schölkopf. Learning with hypergraphs:
529 Clustering, classification, and embedding. *Advances in neural information processing systems*,
530 19, 2006.

- 531 32. Yue Gao, Zizhao Zhang, Haojie Lin, Xibin Zhao, Shaoyi Du, and Changqing Zou. Hypergraph
532 learning: Methods and practices. *IEEE Transactions on Pattern Analysis and Machine*
533 *Intelligence*, 44(5):2548–2566, 2020.
- 534 33. Yifan Feng, Haoxuan You, Zizhao Zhang, Rongrong Ji, and Yue Gao. Hypergraph neural
535 networks, 2019.
- 536 34. Muhan Zhang, Zhicheng Cui, Shali Jiang, and Yixin Chen. Beyond link prediction:
537 Predicting hyperlinks in adjacency space. In *Proceedings of the AAAI conference on artificial*
538 *intelligence*, volume 32, 2018.
- 539 35. Govind Sharma, Prasanna Patil, and M Narasimha Murty. C3mm: Clique-closure based
540 hyperlink prediction. In *IJCAI*, volume 20, pages 3364–3370, 2020.
- 541 36. Naganand Yadati, Vikram Nitin, Madhav Nimishakavi, Prateek Yadav, Anand Louis, and
542 Partha Talukdar. Nhp: Neural hypergraph link prediction. In *Proceedings of the 29th ACM*
543 *international conference on information & knowledge management*, pages 1705–1714, 2020.
- 544 37. Philippe Schwaller, Daniel Probst, Alain C Vaucher, Vishnu H Nair, David Kreutter, Teodoro
545 Laino, and Jean-Louis Reymond. Mapping the space of chemical reactions using attention-
546 based neural networks. *Nature machine intelligence*, 3(2):144–152, 2021.
- 547 38. Seyone Chithrananda, Gabriel Grand, and Bharath Ramsundar. Chemberta: large-scale self-
548 supervised pretraining for molecular property prediction. *arXiv preprint arXiv:2010.09885*,
549 2020.

- 550 39. A Vaswani. Attention is all you need. *Advances in Neural Information Processing Systems*,
551 2017.
- 552 40. Stefanía Magnúsdóttir, Almut Heinken, Laura Kutt, Dmitry A Ravcheev, Eugen Bauer, Alberto
553 Noronha, Kacy Greenhalgh, Christian Jäger, Joanna Baginska, Paul Wilmes, et al. Generation
554 of genome-scale metabolic reconstructions for 773 members of the human gut microbiota.
555 *Nature biotechnology*, 35(1):81–89, 2017.
- 556 41. David B Bernstein, Snorre Sulheim, Eivind Almaas, and Daniel Segrè. Addressing uncertainty
557 in genome-scale metabolic model reconstruction and analysis. *Genome Biology*, 22:1–22,
558 2021.
- 559 42. Wei Lü, Juan Du, Nikola J Schwarzer, Elke Gerbig-Smentek, Oliver Einsle, and Susana LA
560 Andrade. The formate channel foca exports the products of mixed-acid fermentation.
561 *Proceedings of the National Academy of Sciences*, 109(33):13254–13259, 2012.
- 562 43. Max van ‘t Hof, Omkar S Mohite, Jonathan M Monk, Tilmann Weber, Bernhard O Palsson,
563 and Morten OA Sommer. High-quality genome-scale metabolic network reconstruction of
564 probiotic bacterium escherichia coli nissle 1917. *BMC bioinformatics*, 23(1):566, 2022.
- 565 44. Xiao Bu, Jing-Yuan Lin, Jing Cheng, Dong Yang, Chang-Qing Duan, Mattheos Koffas,
566 and Guo-Liang Yan. Engineering endogenous abc transporter with improving atp supply
567 and membrane flexibility enhances the secretion of β -carotene in saccharomyces cerevisiae.
568 *Biotechnology for Biofuels*, 13:1–14, 2020.

- 569 45. Antoine Danchin. Zinc, an unexpected integrator of metabolism? *Microbial Biotechnology*,
570 13(4):895–898, 2020.
- 571 46. Wei Lü, Juan Du, Nikola J. Schwarzer, Elke Gerbig- Smentek, Oliver Einsle, and A. Andrade,
572 Susana L. The formate channel foca exports the products of mixed-acid fermentation.
573 *Proceedings of the National Academy of Sciences*, 109(33):13254–13259, 2012.
- 574 47. Janna Hastings, Paula De Matos, Adriano Dekker, Marcus Ennis, Bhavana Harsha, Namrata
575 Kale, Venkatesh Muthukrishnan, Gareth Owen, Steve Turner, Mark Williams, et al. The chebi
576 reference database and ontology for biologically relevant chemistry: enhancements for 2013.
577 *Nucleic acids research*, 41(D1):D456–D463, 2012.
- 578 48. Tomas Mikolov. Efficient estimation of word representations in vector space. *arXiv preprint*
579 *arXiv:1301.3781*, 2013.
- 580 49. Yue Gao, Yifan Feng, Shuyi Ji, and Rongrong Ji. Hggn+: General hypergraph neural networks.
581 *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 45(3):3181–3199, 2022.
- 582 50. Jianwen Jiang, Yuxuan Wei, Yifan Feng, Jingxuan Cao, and Yue Gao. Dynamic hypergraph
583 neural networks. In *IJCAI*, pages 2635–2641, 2019.
- 584 51. Sunwoo Kim, Soo Yong Lee, Yue Gao, Alessia Antelmi, Mirko Polato, and Kijung Shin. A
585 survey on hypergraph neural networks: An in-depth and step-by-step guide. In *Proceedings of*
586 *the 30th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, pages 6534–
587 6544, 2024.

- 588 52. Jaehyuk Yi and Jinkyoo Park. Hypergraph convolutional recurrent neural network. In
589 *Proceedings of the 26th ACM SIGKDD international conference on knowledge discovery &*
590 *data mining*, pages 3366–3376, 2020.

591 **Competing interests**

592 The authors declare that they have no competing interests.

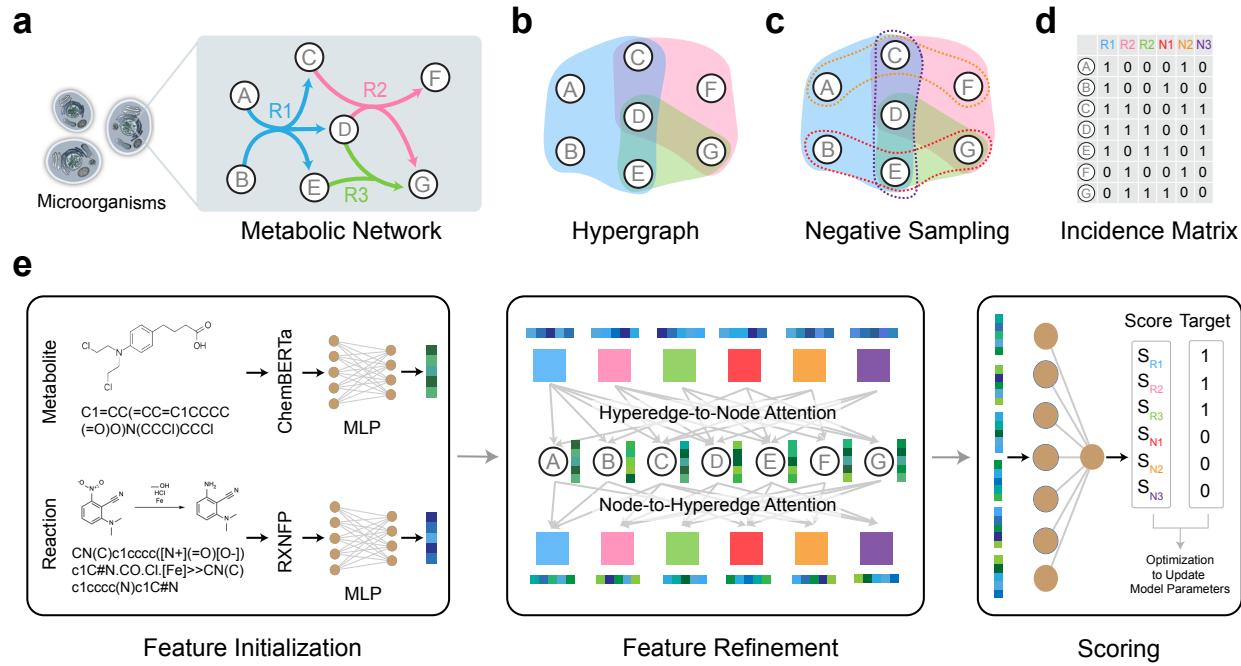


Fig. 1: **MuSHIN workflow.** **a.** Schematic representation of a metabolic network. Metabolites (A-G) are nodes, and reactions (R1, R2, R3) are directed edges connecting reactants and products. **b.** Hypergraph representation of the metabolic network. Each hyperlink connects multiple metabolites participating in the same reaction. **c.** Negative sampling generates chemically invalid reactions (N1, N2, N3) by perturbing positive reactions (R1, R2, R3). Solid and dashed outlines represent positive and negative reactions, respectively. **d.** Incidence matrix of the hypergraph, encoding the binary relationships between metabolites (A-G) and reactions (R1, R2, N1, etc.). **e.** Model pipeline including: (1) feature initialization: ChemBERTa and RXNFP embed metabolites and reactions from SMILES representations, processed via MLPs to extract chemical features; (2) feature refinement: a hypergraph neural network refines node and hyperedge features through alternating hyperedge-to-node and node-to-hyperedge attention mechanisms; (3) scoring: refined reaction features are used to predict confidence scores, compared to target labels for model optimization.

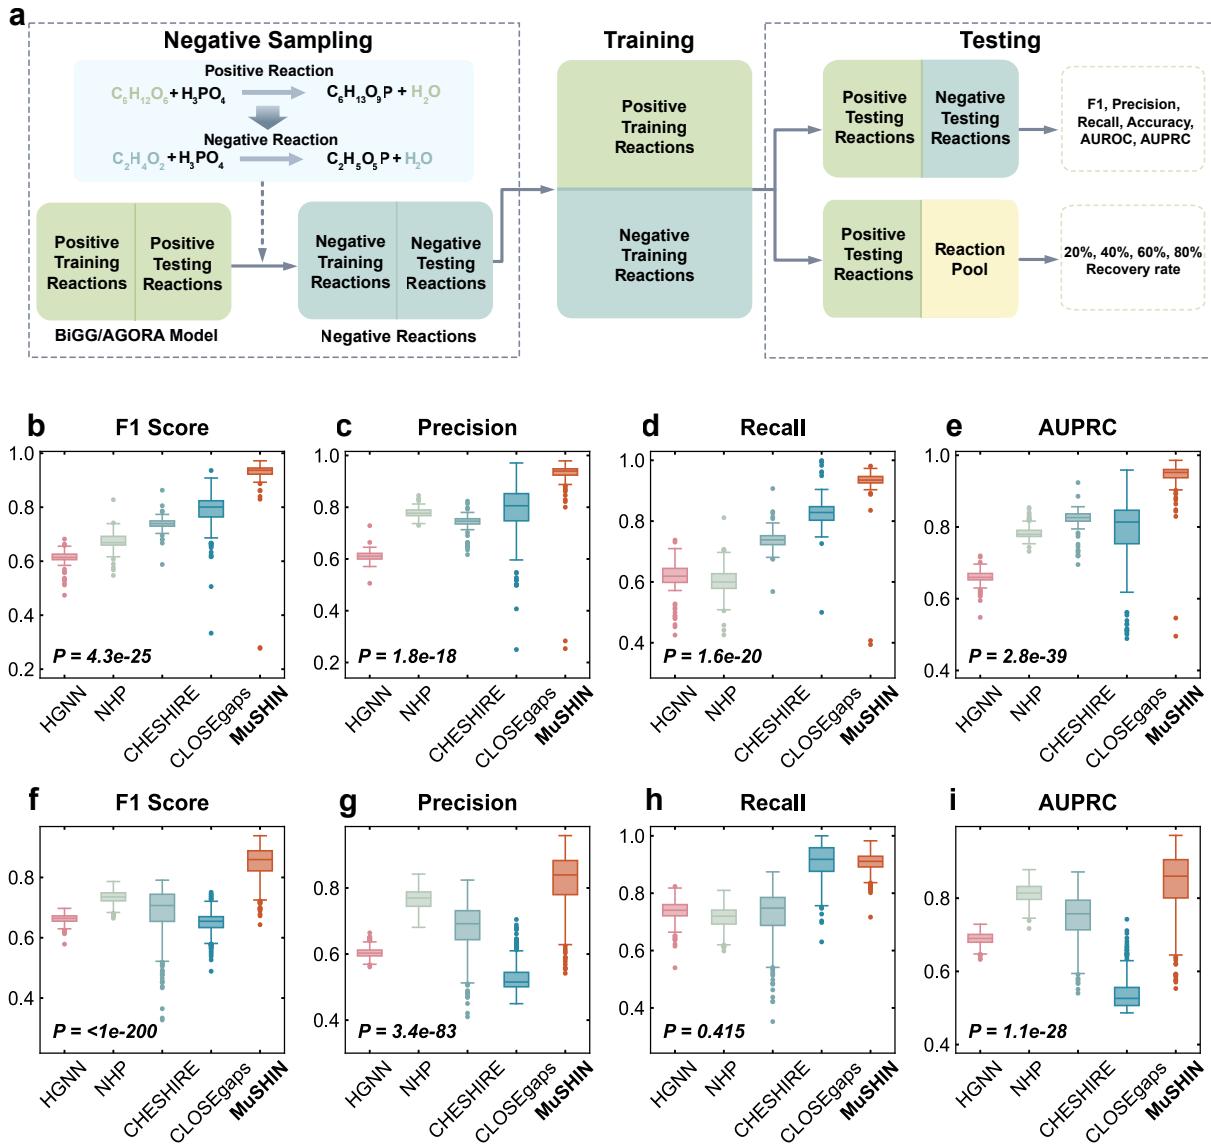


Fig. 2: Overview and internal validation on BiGG and AGORA GEMs. **a.** Workflow of internal validation. Positive reactions from BiGG/AGORA GEMs are split into training and testing sets, while negative reactions are generated by perturbing valid reactions to create chemically invalid ones. Models are trained to distinguish positives from negatives and evaluated using statistical metrics. **b–e.** Internal validation results on 108 BiGG GEMs, evaluated using six metrics (F1 Score, Precision, Recall, AUPRC). Each dot represents one GEM, averaged over 10 Monte Carlo runs. MuSHIN consistently outperforms HGNN, NHP, CHESHIRE, and CLOSEgaps. **f–i.** Internal validation on 818 AGORA GEMs of gut bacteria using the same evaluation setup and metrics. Two-sided paired-sample t-tests were conducted between MuSHIN and the second-best baseline; exact p-values are reported. Each boxplot shows the median (central line), interquartile range (boxes), and variability across GEMs (whiskers and individual points).

Source data are provided as a Source Data file.

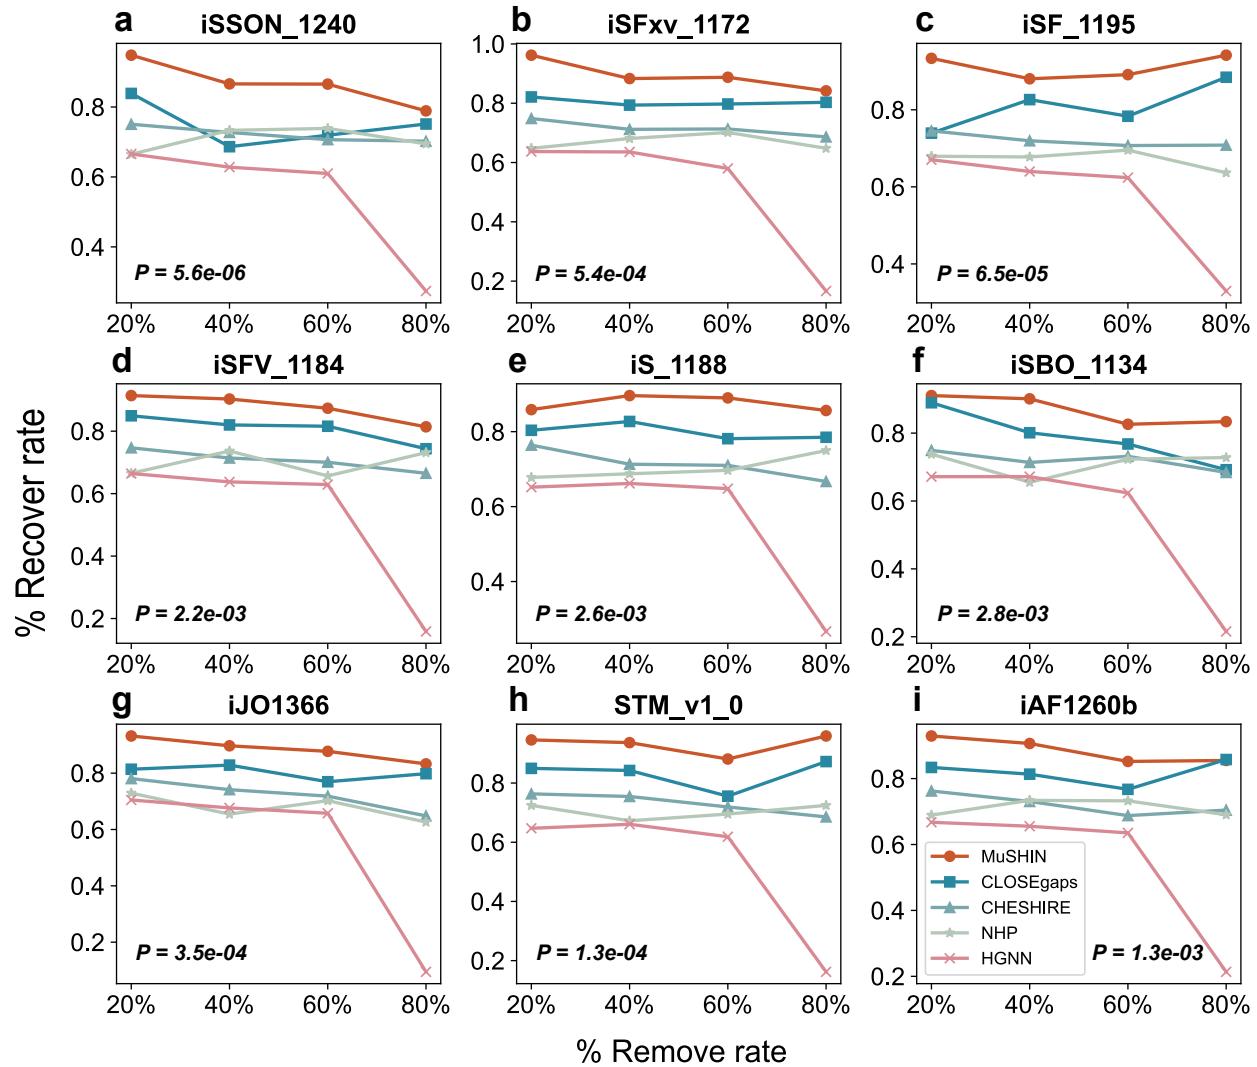


Fig. 3: Recovery performance on highly incomplete GEMs. The recovery rates of MuSHIN and four baseline methods (HGNN, NHP, CHESHIRE, and CLOSEgaps) are evaluated across nine GEMs (iAF1260b, iJO1366, iS_1188, iSBO_1134, iSF_1195, iSFV_1184, iSFxv_1172, iSSON_1240, STM_v1_0) with varying levels of incompleteness (20%, 40%, 60%, and 80% of reactions removed). MuSHIN consistently demonstrates higher recovery rates compared to the baselines across all GEMs and removal rates, with particularly stable performance even under extreme scenarios (80% removal). Each data point represents the average recovery rate from 10 Monte Carlo simulations, while the lines illustrate trends across removal rates. Two-sided independent t-tests were performed between MuSHIN and CLOSEgaps for each model; exact p-values are shown.

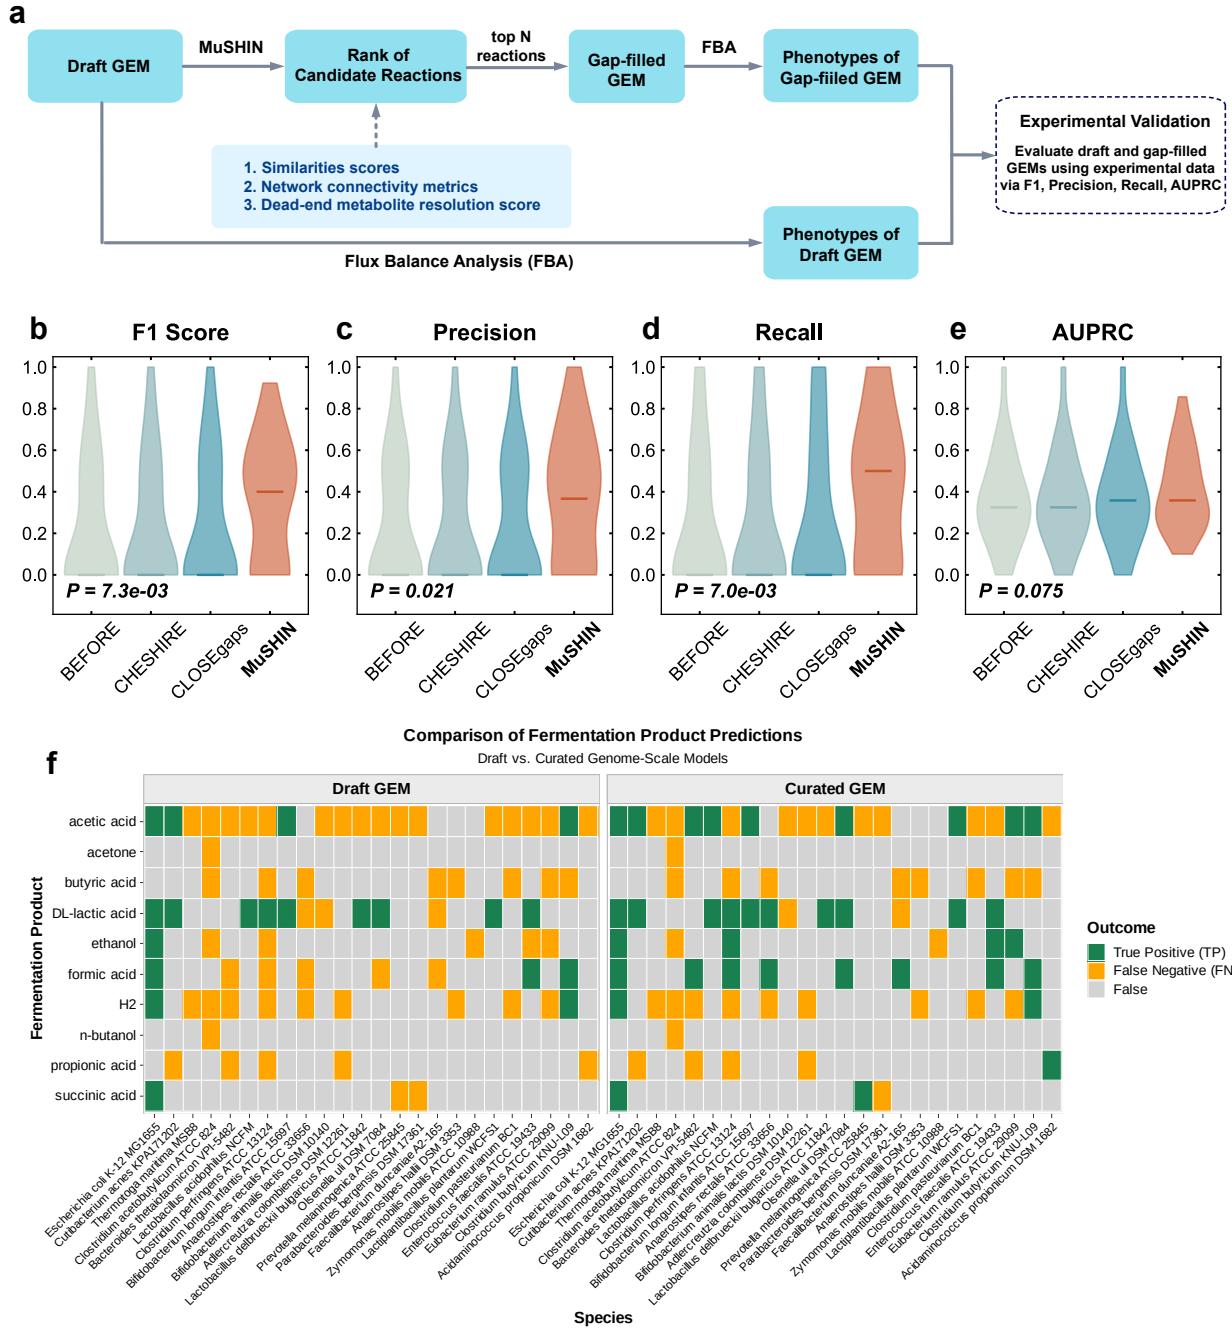


Fig. 4: Overview and external validation performance. a. Workflow of external validation.

MuSHIN ranks candidate reactions for draft GEMs, and top predictions are used to construct gap-filled GEMs. Flux balance analysis is applied to both draft and gap-filled models to simulate phenotypes, which are then compared with experimental data using classification metrics. **b–e.** External validation performance with 100 added reactions on anaerobic bacteria GEMs. Violin plots show F1 Score, Precision, Recall, and AUPRC for MuSHIN, CHESHIRE, CLOSEgaps, and the initial model ('BEFORE'). MuSHIN outperforms CHESHIRE and the initial state, and is competitive or superior to CLOSEgaps across all metrics. Each point represents the mean result from 10 Monte Carlo simulations. Two-sided Wilcoxon signed-rank tests were conducted between MuSHIN and the second-best baseline method; exact p-values are shown in each panel. Each boxplot shows the median (central line), interquartile range (boxes), and variability across GEMs (whiskers and individual points). **f.** Validation of predicted fermentation products. Heatmap compares MuSHIN-predicted fermentation capabilities with experimentally known outcomes across 24 bacterial strains and 10 common fermentation products. Green cells indicate known products successfully recovered by MuSHIN, while orange boxes mark false-positive predictions introduced during gap-filling.